Team B2: Jamming Attack on Voice Activated Systems

Final Presentation

Cyrus Daruwala, Eugene Luo, Spencer Yu

Recap

Motivation: Previous research has shown attacks to covertly manipulate voice activated systems are possible, and we are aiming to show the relevance of these attacks by demonstrating them on commodity hardware.

Goal: reduce the accuracy with which the wake word (“Hey Siri”) is recognized.

Solution:

1. Determine the input to jam the wake word (done)
2. Ensure that the jamming input is within latency bounds (done)
3. Reduce the % of false positives without compromising on the % of false negatives (goal for next week)

Changes to Solution Approach: MFCC + DTW

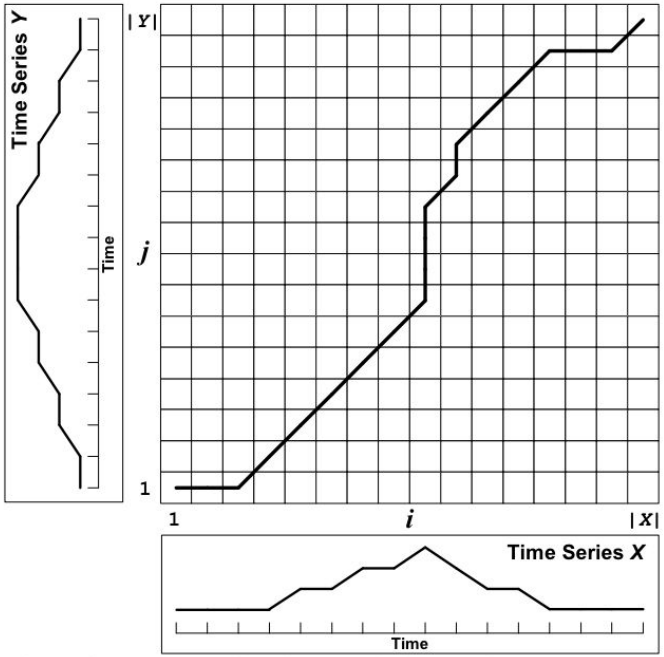

Figure 2. A cost matrix with the minimum-distance warp path traced through it.

Source: https://cs.fit.edu/~pkc/papers/tdm04.pdf

Overall goal: Compare the similarity between two sounds of different lengths by aligning their identifying characteristics.

MFCC: coefficients that characterize the speech

DTW: matches characteristics in sound by stretching/compressing

System Diagram

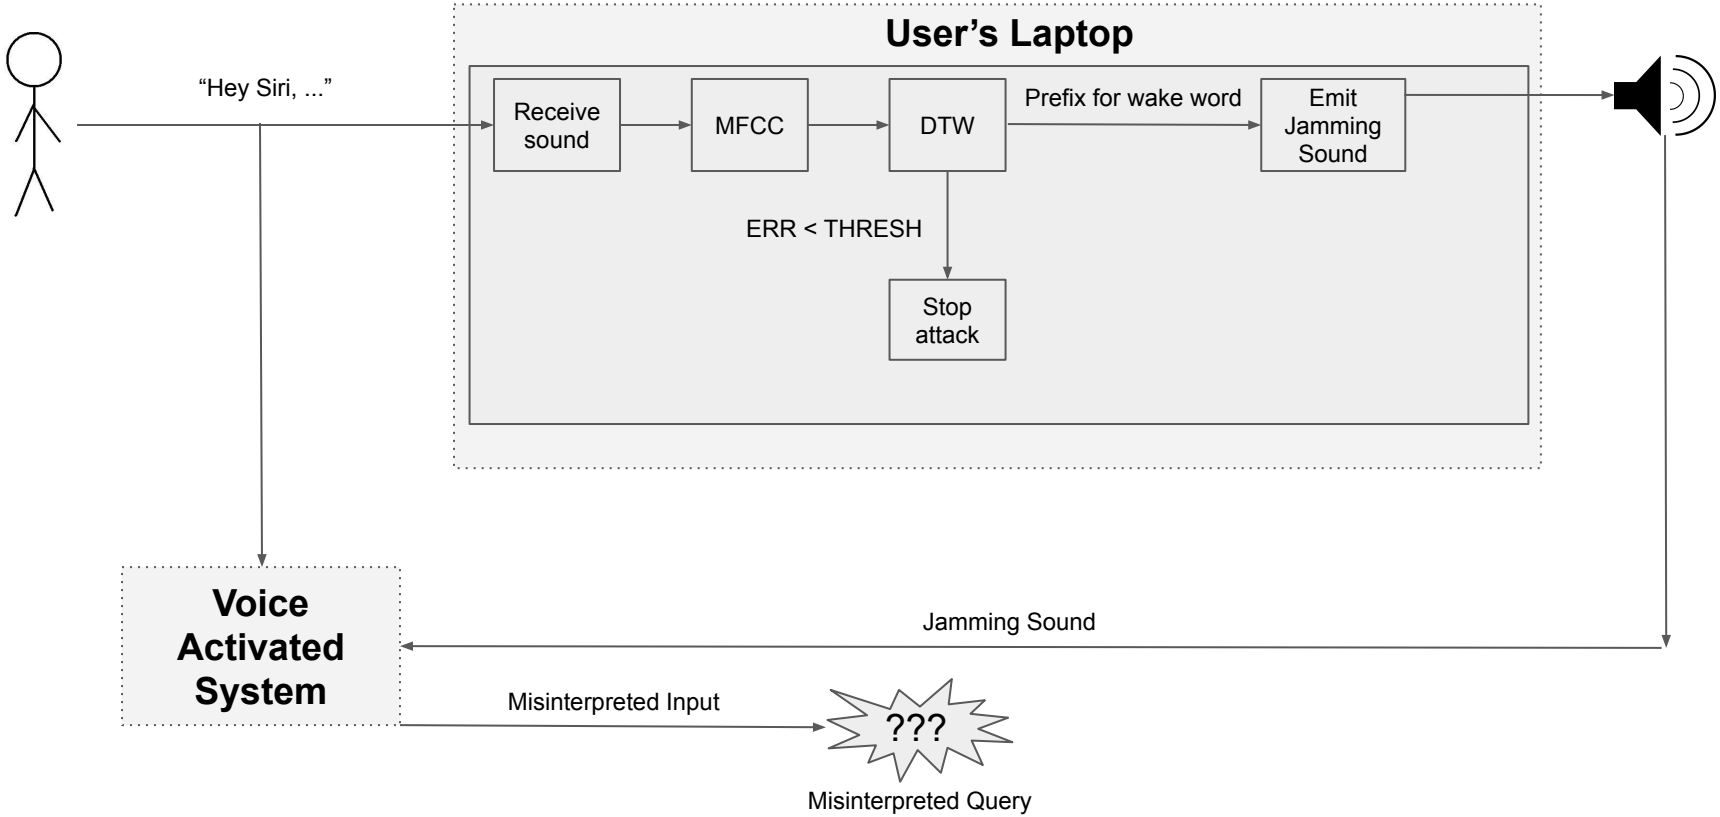

Demo: Complete Solution

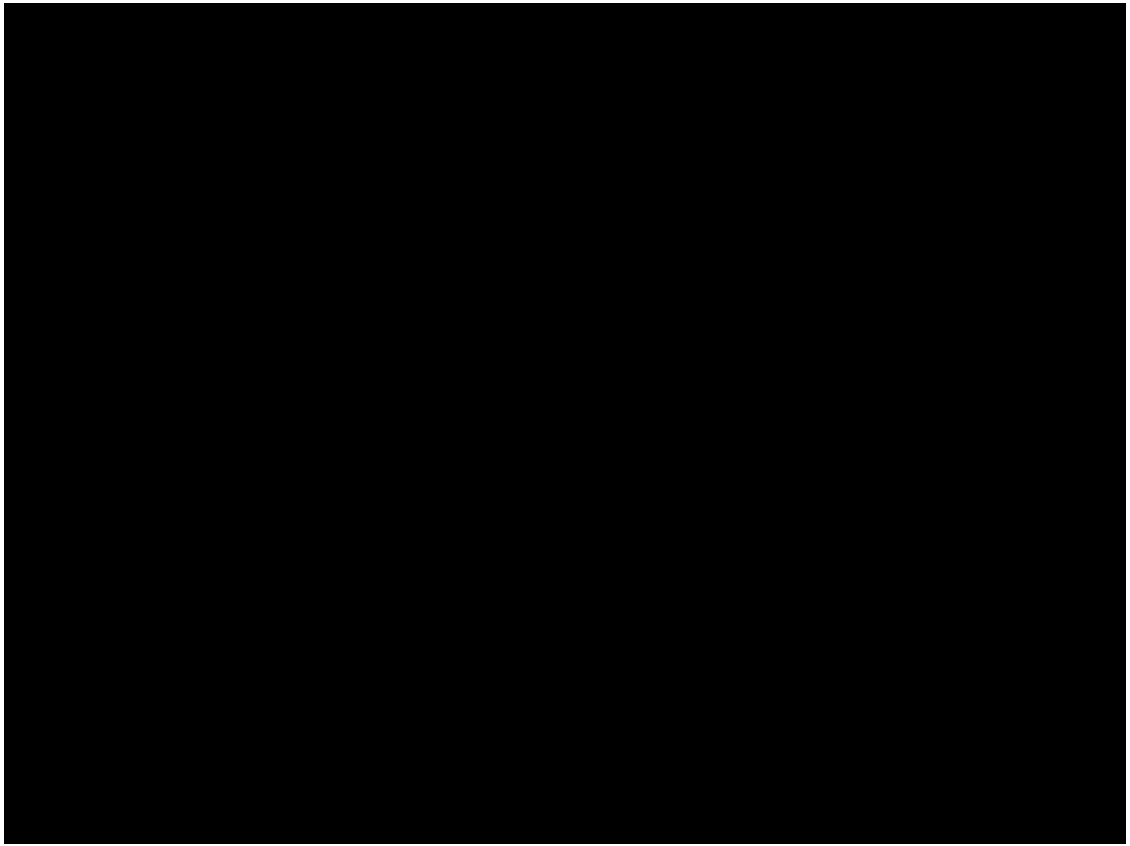

IMG_2782.mov

Demo: Selective Jamming

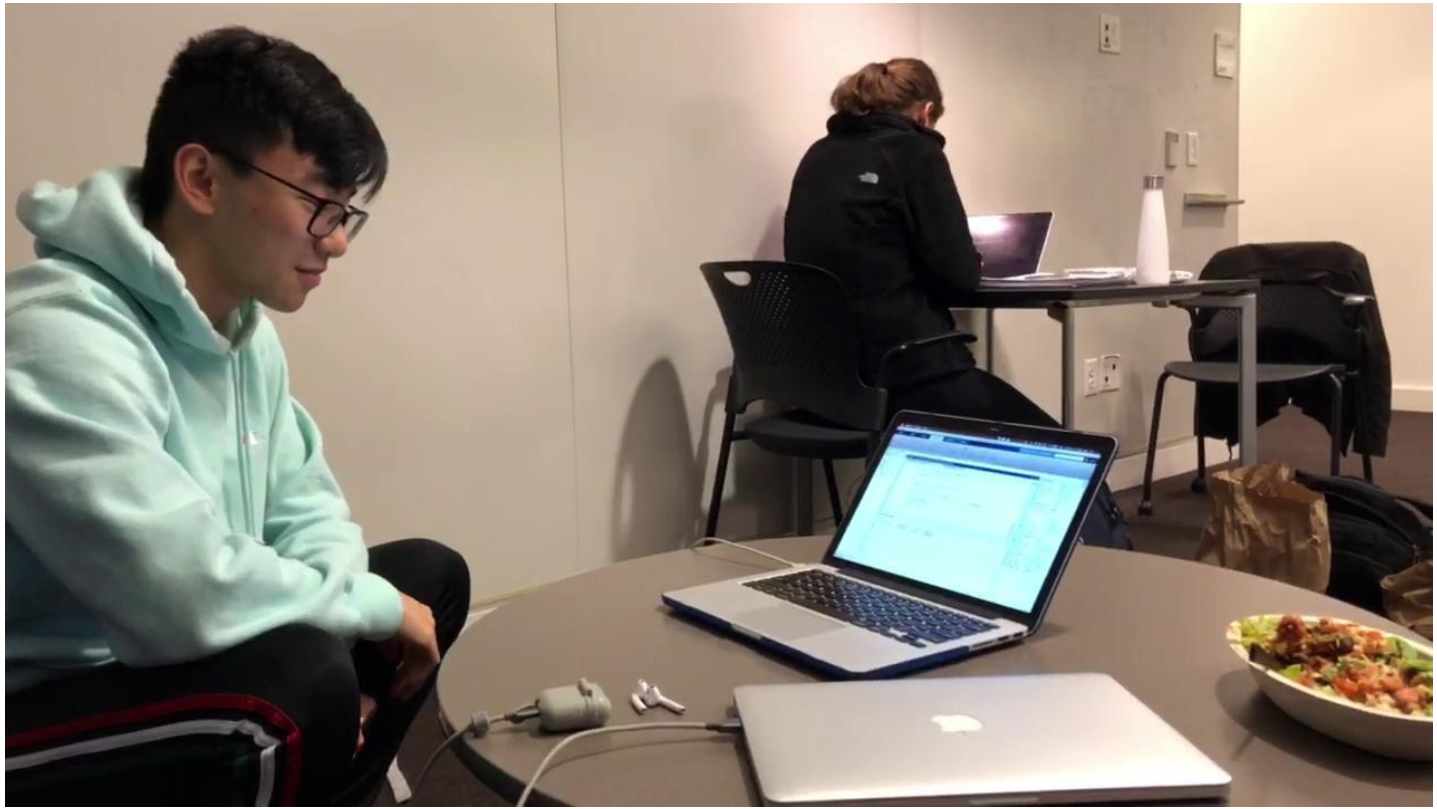

IMG_2786.mov

Measuring Performance

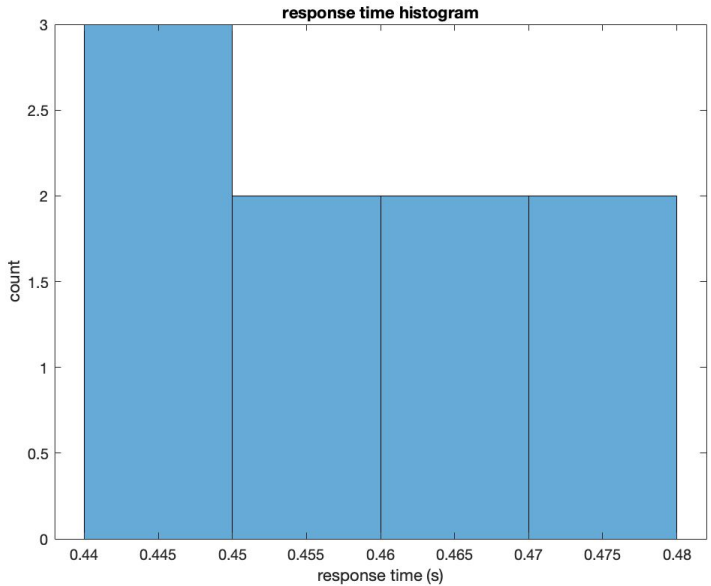

How do we verify the overall **performance and speed** of our system?

Latency: time sync + manual validation via MATLAB scripts

- Response time average **462.5ms** over 10 queries

System Validation: tested against Spencer's iPhone 7 Plus

- Response rate drops to **46.67%** over 30 queries

Metrics by Sound

Noise Type	White Noise	Music w/ vocals[1]	Music w/o vocals[2]	Podcasts[3]	Human Speech[4]
False Negative Rate	N/A	N/A	N/A	0%	0%
False Positive Rate	0%	3.43% 20.3/min	2.105% 15/min	2.7% 17.5/min	0.263%, 5.22/min

False Negative Rate: % where attack does not trigger when it should

False Positive Rate: % where attack triggers when it should not

[1] - Sound of a box fan (https://www.youtube.com/watch?v=Qq_wJFvhfrg)
[1] - “Hey Jude” by the Beatles
[2] - “The Sleeping Beauty, Op. 66, TH 13: No. 6, Valse” by Pyotr Ilyich Tchaikovsky, conducted by Eugene Ormandy
[3] - “Great Bitter Lake Association”, *99% Invisible* by Roman Mars
[4] - Spencer Yu

Validating Signal Processing

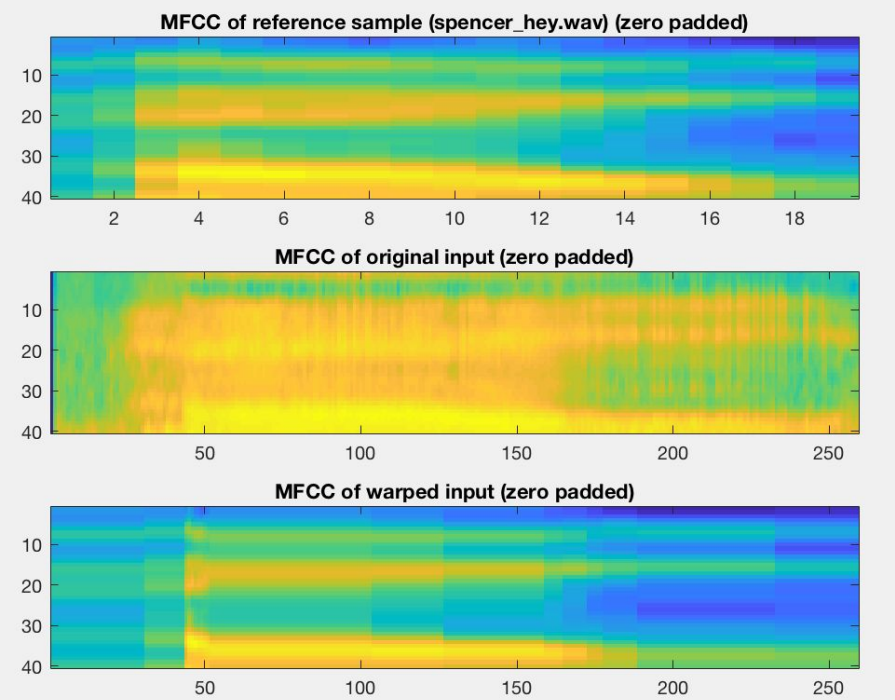

How do we verify each **individual component** of our attack works as expected?

MFCC: IDCT + spectrogram

DTW: IDCT + spectrogram

Play out sounds

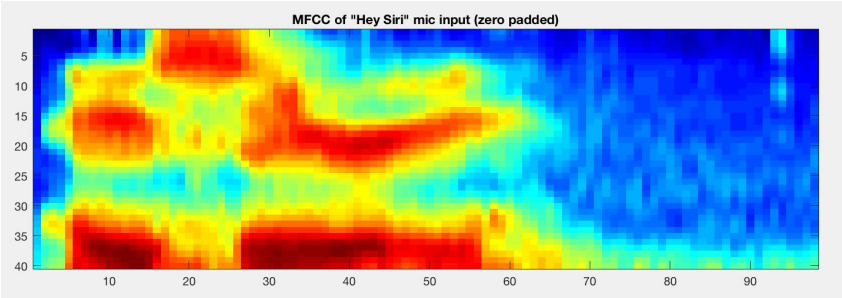

Project Management

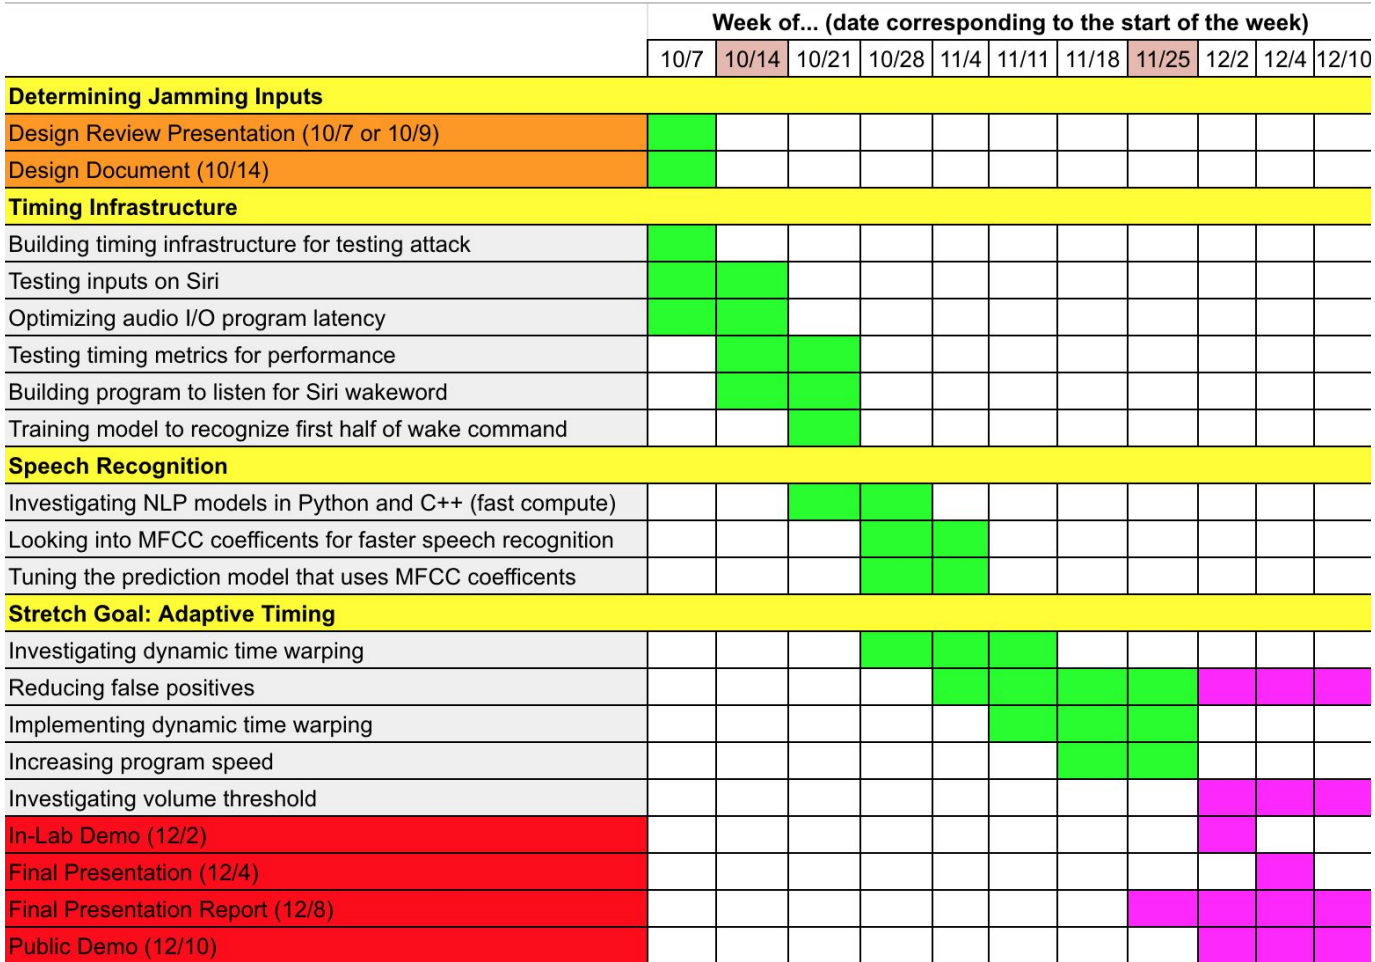

Obstacles Encountered

Choice of language

- **Before: Python**
- **After: Matlab**

Area of ECE Explored in Project

- **Before:** Software with signal support
- **After:** Largely signal processing with software implementation

Solution Approach

- **Before:** ML for generating/triggering adversarial input
- **After:** Signal processing

Lessons Learned

Scope of project

Having backup plans

Constant feedback from experts

Choose interesting projects

Take risks and try something new